

# Chapter 17 Free Space Management.

---

- Managing free space can be easy:
  - When using paging.
  - When the space is divided into fixed-size chunks
    - Use an array listing the chunks.
- When free space managing becomes more complicated.
  - When the free space is of varying sizes.
    - Happens in user level allocations.
    - Happens when OS uses segmentation.
  - External Fragmentation:
    - The free space gets chopped up into little pieces of different sizes.
    - Subsequent requests for memory may fail because there is no single contiguous space of large enough size.

## 17.1 Assumptions

- Assume a basic interface that returns a pointer to a space of a requested size in memory.
- When freeing the memory the library must know the size of memory to be freed.
  - The library must know how big a chunk of memory is just by being handed a pointer to it
- The generic data structure used to manage the free space on the heap is a free list.
  - Contains references to all the free chunks of space in the managed region of memory.
- Primarily concerned with external fragmentation.
- Internal fragmentation is when the allocator hands out chunks of

memory larger than requested.

- Unasked for and unused memory is considered internal fragmentation.
- It's internal to the allocation rather than external.
- Assume that once memory is handed out it cannot be relocated to another location in memory.
  - Eliminates the ability of compaction.
- Assume the allocator manages a contiguous region of bytes.
  - In some cases allowing for expansion.

## 17.2 Low-level Mechanisms

- **Splitting and Coalescing**

- A free list contains a set of elements that describe the free space remaining in the heap.
- Assume there is a request for a single byte of memory:
  - The allocator will perform splitting.
    - The allocator will find a free chunk of memory that can satisfy the request and split it in two.
    - The first chunk will return to the caller.
    - The second will remain on the free list.
  - The list essentially stays intact just missing the allocated portions.
- Coalescing of free space.
  - If memory is freed returning space in the middle of the heap.
  - If not careful it may divide contiguous spaces into multiple segments.
    - Could have memory request fail because the space is divided into chunks that are too small.
  - If the newly freed space is next to previously freed space

then it is combined into one chunk.

- **Tracking the Size of Allocated Regions**

- It is assumed that given a pointer the malloc library can determine the size of a region of memory.
- Most allocators store a bit of extra information in the header block which is kept in memory just before the handed out chunk of memory.
  - The header minimally contains the size of the allocated region.
  - May contain additional pointers to speed up deallocation.
  - A magic number to provide additional integrity checking.
    - Used as a sanity check.

- **Embedding a Free List**

- How do we build a list inside the free space?
  - Can't use malloc because it's being done inside the malloc library.
- First the list has to be initialized.
- The heap is built calling mmap().
- Then splitting and coalescing is used to manage the memory.

- **Growing the Heap**

- What should be done when the heap runs out of space?
  - The simplest approach is to let it fail.
  - Most traditional allocators start with a small heap and request more memory from the OS when it runs out.

## 17.3 Basic Strategies

- The ideal allocator is both fast and minimizes fragmentation.

- **Best Fit**

- Search through the free list and find chunks of free memory that are as big or bigger than the requested size.

- Return the one that is the smallest in the group.
- Performance penalty for exhaustive block search.
- **Worst Fit**
  - Find the largest chunk and return the requested amount.
  - Still exhaustive block search.
  - Mostly shown to perform badly resulting in excess fragmentation.
- **First Fit**
  - Finds the first block that is big enough and returns requested amount.
  - Fast method.
  - Can pollute the beginning of the free list with small objects.
- **Next Fit**
  - Keeps a pointer to the location within the list where one was looking at last.
  - Performance is similar to First Fit.
  - Designed to prevent the fragmentation of first fit.

## 17.4 Other Approaches

- **Segregated Lists**
  - If an application has one popular sized request that it makes keep a separate list just to manage that size object.
    - All other requests are forwarded to the main list.
  - Complications:
    - How much memory should be dedicated to the pool.
    - The slab allocator handles that problem.
      - at boot time it allocates a number of object caches for kernel objects likely to be requested frequently.
      - When a cache is running low on memory it

requests a slab of memory from the free list.

- When the reference counts for a slab are zero the OS can reclaim the unused slab.
- The slab allocator keeps the free objects on the list in a pre-initialized state.

- **Buddy Allocation**

- Allocator designed around making coalescing simple.
- Memory is first thought of as one big space of size  $2^N$ .
- When memory is requested a recursive search splitting space available in halves until the proper size block is found.
- Scheme can suffer from internal fragmentation.
  - Blocks are only allowed to be allocated in sizes of powers of 2.
- Coalesces at the same time it frees memory.
  - checking is easy because it checks addresses that differ by single bits.

- **Other Ideas**

- The ideas above suffer from a lack of scaling.
  - List searches can be slow so more advanced data structures are used.
    - Balanced binary trees, Splay trees, partially ordered trees, etc.